

Take A Stand: Fight Moral Grandstanding on Twitter

I. The Problem

The problem I would like to address is the prevalence of moral grandstanding on social media platforms, exacerbated by the gamification of communication. For the purpose of this paper, I will narrow the scope of social media down to Twitter; however, most of the statements and solutions regarding Twitter can be generalized to other addictive social media platforms.

Moral grandstanding is a form of pretense to make oneself appear more virtuous in the eyes of others. It is difficult to distinguish the genuinely concerned from the moral grandstander; the difference lies in the intent of the speaker. As Tosi and Warmke's Aeon article articulates, "Less sincere grandstanders might not care one way or another about their stated cause, but still want others to *believe* that they care." Thus, while the grandstander may even be speaking the truth, the primary goal is to represent themselves on higher moral ground than everyone else.

Explanations for moral grandstanding are largely psychological. Tosi and Warmke reference psychological studies that found that not only do we find ourselves morally superior to others, we also want others to think that about us too. In addition, moral grandstanding could be seen as a form of socially adaptive behavior: we present socially respected moral statements to be accepted better by our community. On social media, the entire structure of the platform reinforces the innate psychological drive to grandstand. Thi Nyugen's analysis of the gamification of Twitter highlights how the system of Likes and Retweets is ideal for grandstanding, because the entire goal of the game is to win external approval and validation through quantifiable measures. Unless the user is a value-independent gameplayer who uses the Likes and Retweets purely as measurement tools, anyone who is playing the game, intentionally or unintentionally, will find that grandstanding is one of the easiest paths towards gaining Likes and Retweets and winning popularity and power on Twitter.

II. The Stakes

Some dangerous effects of moral grandstanding are:

1. Cynicism: because it is so difficult to distinguish between grandstanding and true moral outrage, people may begin to doubt the trustworthiness of everyone's claims, leading to a derisive outlook towards moral discussion as a whole.

2. Polarization: if two people are part of different groups with different moral opinions on a matter, partaking in grandstanding within their own communities will increase the divergence between the two groups.

3. Silencing: moderate voices are no longer heard because they do not attract attention. Particularly on social media, polarization is especially easy because one can choose who they interact with and who they block, and silencing is overlooked because there are so many voices to replace the silenced voice.

The ultimate danger of these effects is that people will disengage. John Stuart Mill argues that “the truth depends on a balance to be struck between two sets of conflicting reasons.” If people no longer see the purpose of discussion, people will stop engaging in the healthy conflict of opinions that Mill encourages – and without this conflict, Mill argues, it is much more difficult for society to discern truth.

III. Barriers to Intervention

What makes moral grandstanding on Twitter so difficult to address? First, grandstanding on its own is already hard to reduce, as it is a natural tendency for people to seek external validation, attained through presenting a “more virtuous” self image. In addition, as Tosi and Warmke state, we can’t solve grandstanding by accusing people who seem like they are, because we don’t actually know their intentions. Instead, reducing grandstanding has to come from an internal desire to change one’s intention. Psychological reboots cannot be enforced by external policy changes, which makes the elimination of grandstanding difficult.

In addition, Twitter rewards people who grandstand by displaying Likes and Retweets, which encourage people to prioritize pleasing their community over presenting genuine thoughts. If we remove the gamified incentives then, the users would be less likely to engage in grandstanding, as it would be more difficult to tell if they are amassing positive public opinion. However, without gamification, the participation rates on platforms like Twitter would drop significantly. The goal, then, is to shift gamification in a way that promotes genuine communication over moral grandstanding and extremism.

IV. A Potential Solution

My proposed solution is to modify the gamification of Twitter as well as promote self-reflection.

We begin with the first approach: modifying gamification to promote deeper engagement and conversation. Two immediate changes will lay the groundwork for more detailed

modifications: 1. Removing the word limit to stop encouraging brief and extreme statements, and 2. Creating a “social engagement” metric as a new measure of “success” on Twitter. Subsequently, the main changes for reducing grandstanding can be made:

Twitter can implement two new buttons: a “Changed My Mind” button, which shows that the post changed the reader’s mind, and an “Intrigued” button, which indicates that the reader may not agree but finds the argument interesting. Both new buttons would add points to one’s social engagement rank.

Additionally, Twitter should encourage intercommunity discussion by tracking the amount of intercommunity discourse inspired by a post – responses from non-followers are weighted more than responses from followers. This reduces the urge to grandstand because diversity of opinion clouds the position that “should” be taken by the grandstander.

Both of these changes reduce grandstanding because knowing that one’s post reaches a wide audience with diverse opinions blurs the clarity of which stance to take when grandstanding and prompts more moderately phrased posts to appeal to people with differing opinions.

Addressing Objections:

1. People write bland posts to appeal to the largest audience possible.

However, in order to get people to change their mind or find the post intriguing, there has to be some bite to the post; otherwise, it would be entirely overlooked. Thus, the game is now a puzzle of how to engage with the most people without becoming opinionless. This is also why there is no minimum word limit – if a person is able to communicate effectively in a sentence, then they should not be forced to write more than they need to.

2. People in echo chambers may sink deeper into their echo chamber as a result of being exposed to voices outside of their community.

This brings up the possibility that other problems of the post-truth age may worsen as a side effect. In order to bring people out of echo chambers, Thi Nyugen suggests that we must “work to repair the broken trust between echo chamber members and the outside social world.” I would argue that no social media platform is ideal for repairing broken trust, due to their impersonal nature. However, these changes may build some foundations for trust. With modern technology, people form online friendships all the time, so it seems to be possible to build trust virtually. While mere exposure won't break the echo chamber, an outsider, either out of kindness or incentivized by the “Changed My Mind” button, may be able to befriend the echo chamber member through intercommunity responses. Furthermore, while these changes may not be able to truly induce the social epistemic reboot necessary to break echo chambers, they will be able to draw people out of epistemic bubbles through exposure to information beyond their follower community, which shows that it addresses another problem of the post-truth era beyond grandstanding.

3. People stop using Twitter because non-followers are wrecking their community. This is a real concern, because loss of engagement is the last thing we want to promote. We can mediate this by putting non-follower responses in a separate page that people click into – the most viewed (by followers and non-followers) responses will begin to populate the main thread as well. This prevents the main thread from being inundated with negative responses, while still presenting a fair share of external opinions.

Ideally, the gamification is still enough to encourage people to use Twitter for good communication purposes. This gamification is still fallible to value capture as people who engage would not necessarily be chasing genuine communication but rather are still chasing the tangible measure of higher social engagement rank. However, this form of value capture still results in overall reduced grandstanding and polarization.

We now turn to the second approach that will ideally shape users into more reflective posters. In the words of Tosi and Warmke, “thinking about grandstanding is a cause for self-reflection , not a call to arms. ... Are we doing good with our moral talk? Or are we trying to convince others that we are good?” This is precisely the question that the user must answer internally. However, it is too easy right now to quickly write down the first thoughts that come to mind and send it out into the world without even reading it over again, much less think about the impact. Thus, a set of brief questions would force the user to pause. The questions could look like this:

1. Why are you posting this?
2. Are you happy with that intention?
3. Who is your target audience?
4. How does this contribute meaningfully to the discussion?

Addressing Objections:

1. People will not write thoughtful responses.

The system is evidently easy to circumvent, as the poster can lie or refuse to write meaningful sentences; however, it still induces a moment of reflection that is especially effective for people who are not aware of their grandstanding. Even for people who are willing to continue grandstanding while fully knowing what they are doing, these questions can perhaps inspire a little bit of guilt or make them rethink their actions.

2. People will find this annoying and stop using Twitter.

One way to prevent this is to use a tone detector. We already have language-processing AI that can detect inflammatory language and tone; in addition, we can flag responses that are too short – for example, less than five words. We will present these questions only when a response is flagged – this way, most thoughtful posts will not undergo the interrogation.

Furthermore, if the poster genuinely desires to post an inflammatory post for a cause they believe in, they will be willing to answer these questions; thus, it also filters for genuinity. The responses to these questions will be posted publicly beside the actual post, allowing readers to decide for themselves whether this post is worth reading.

To summarize, my proposal is to change Twitter's platform in two ways: 1. Gamify the platform to encourage genuine discussion over moral grandstanding within one's closed follower community, and 2. Create self-reflection through questions that interrogate one's intent. The first addresses the human tendency to respond to gamification. The second shapes the values of the users to become value independent game players. With funding from the nonprofit, we can provide financial incentive for Twitter to implement these changes and reduce moral grandstanding on the platform.